

APA- Practical Work 2017–2018

Albert Ribes Kerstin Winter

January 17, 2018

Contents

1	Introduction	3
1.1	Description of the work and its goals	3
1.2	Description of available data	3
1.3	Needed packages	3
2	Related Previous Work	3
3	Data exploration process	4
3.1	Preprocessing	4
3.1.1	Treatment of missing values	4
3.1.2	Treatment of anomalous values	4
3.1.3	Treatment of incoherent values	4
3.1.4	Coding of non-continuous or non-ordered variables	4
3.1.5	Possible elimination of irrelevant variables	4
3.1.6	Creation of new useful variables (Feature extraction)	5
3.1.7	Normalization of the variables	5
3.1.8	Transformation of the variables	5
3.2	Renaming of the attributes	5
4	Resampling protocol	6
5	Metric used to evaluate the models	7
6	Results obtained using linear/quadratic methods	7
6.1	Naive Bayes	7
6.2	K-nearest Neighbours	8
6.3	General Linear Model	9
7	Results obtained using non-linear methods	10
7.1	Random Forest	10
7.2	Neural Network	10
8	Description and justification of the final model chosen and estimation of the generalization error	11

9	Self-assessment of successes, failures and doubts	12
9.1	Successes	12
9.2	Failures	12
9.3	Doubts	12
10	Scientific and personal conclusions	12
11	Possible extensions and known limitations	12

1 Introduction

1.1 Description of the work and its goals

The goal of this project is to build a classification model to predict whether a lung cancer patient will die within one year after surgery or not. To do so we will study a dataset with real lung cancer patients.

As this is very sensitive information, our priority will be to minimize the amount of false negatives, i.e., avoid predicting a patient will not die within one year when he actually does.

The data is taken from [2]

1.2 Description of available data

The data we are working with is about patients who underwent major lung resections for primary lung cancer in the years from 2007 to 2011. For each patient we are given information about his diagnosis and effects produced by the cancer.

The dataset is very limited in the number of instances available as it only contains data about 470 patients. In addition, the distribution of the predicted class is not quite balanced, since only 70 of the patients died in a one year period after their surgery. This may become a problem in some of the prediction models due to the fact that the results will be biased towards the bigger class. However, we can suppose that the data has been collected uniformly and that this proportion is similar to the probability of the real population.

For each patient we have 16 different attributes of which three are numerical, and the rest is categorical. From those, ten are binary variables. The response attribute is also binary.

1.3 Needed packages

The following R packages have been used in this work:

- klaR
- randomForest
- caret
- ROCR
- kknn

2 Related Previous Work

This is a very tricky dataset due to the imbalance of the data and the little amount of instances given. As this is a very typical situation in daily normal studies regarding Machine Learning, many groups of researchers are involved on predicting data of this kind of situations.

The paper [2] suggests using some special ensambling methods together with cost sensitive Support Vector Machines to solve imbalance data problems.

3 Data exploration process

3.1 Preprocessing

3.1.1 Treatment of missing values

Our dataset does not have missing values, so there is no need to treat them.

3.1.2 Treatment of anomalous values

The age of the patients is not well distributed. Most of them are over 60 years old, and only four of the patients are under 40. Due to this, it is very likely that the conclusions of our study will only be applicable to the elder people. However, we are reluctant to remove the younger patients.

3.1.3 Treatment of incoherent values

The variable FEV1, which is the Forced Expiratory Volume in 1 second, shows a few anomalously high values. Depending on factors like age and sex of a patient the average value of the FEV1 should be around 3-6 litres, whereas the dataset shows values up to 86. As most of the values are within 0 and 10 we assume that the dataset contains the FEV1 in litres. To decide which values are to be determined as outliers we calculate the FEV1/FVC ratio which gives the percentage of the lung volume exhaled in the first second over the whole exhaled volume. All patients having a unrealistic ratio higher than 100%, which are 22 patients, are determined to be outliers and eliminated. We chose not to apply any stricter constraints because the dataset does not include the sex of the patients which influences the normal values of the FEV1 much.

Source of knowledge about FEV1 and FVC: [1]

3.1.4 Coding of non-continuous or non-ordered variables

As most of the dataset variables are logical, we have coded them as logical in R. We have converted the variable “DGN” to many binary variables, each one saying whether the patient showed that diagnosis or not.

Originally, the variables “PERFORMANCE” and “SIZE” were categorical, but as they seem to have some kind of order, we have coded them as numeric. The “PERFORMANCE” can have the values {1,2,3} and the variable “SIZE” can have the values {1,2,3,4}. Both of them are then normalized.

3.1.5 Possible elimination of irrelevant variables

Some of the variables of our dataset are not well represented. In particular:

DGN	There are just one patient with DGN = 1 and 8 with DGN = 8
PAD	Just 8 patients have PAD = True
ASHTMA	Just 2 patients have ASHTMA = True
MI	Just 2 patients have MI = True

As we have very few instances we will train and run the models on two different datasets. One of them, `thoracic.original`, will contain all of the original attributes, and the other, `thoracic.removed`, will contain all attributes but “DGN.1”, “DGN.8”, “PAD”, “ASHTMA” and “MI”.

3.1.6 Creation of new useful variables (Feature extraction)

We have not created any new variable as most of the ones we have are logical and at least without medical knowledge we cannot determine any relation they might have. The only relation we found is the FEV1/FVC ratio, but as we cannot ensure that these do not influence in other variables we did not replace these variables by their ratio.

3.1.7 Normalization of the variables

We need to normalize only numeric variables, which are AGE, FVC and FEV1. For FVC and FEV1 the range will correspond to the maximum and minimum observed values with a margin of 10%. To normalize AGE we will only consider cases between 0 and 100 years old.

As we have converted the variables “PERFORMANCE” and “SIZE” to numeric, these are also normalized within 1-3 and 1-4.

3.1.8 Transformation of the variables

Searching through the web we have found that the acceptable range for skewness in a numeric variable is $(-2, +2)$. In the original dataset, the one which contains all the patients, the skewness of the variables AGE, FVC and FEV were out of this range. But after eliminating 22 of the patients (in 3.1.3) all of them are inside the acceptable range, so there’s no need to transform them.

3.2 Renaming of the attributes

In order to improve legibility we have renamed some of the attributes in the original dataset to better understand the meaning of each one. The following table maps the name in the original dataset to our pre-processed dataset.

Original	Our naming
DGN	DGN
PRE4	FVC
PRE5	FEV1
PRE6	PERFORMANCE
PRE7	PAIN
PRE8	HAEMOPTYSIS
PRE9	DISPNOEA
PRE10	COUGH
PRE11	WEAKNESS
PRE14	SIZE
PRE17	DIABETES
PRE19	MI
PRE25	PAD
PRE30	SMOKE
PRE32	ASTHMA
AGE	AGE
Risk1Y	DIED

4 Resampling protocol

After the preprocessing we have two datasets, each of them with 448 instances. The dataset *thoraric.original* has 23 attributes and the dataset *thoraric.removed* has 19 attributes after the eliminations explained in section 3.1.5. For about 85% of the patients the target variable (DIED) equals FALSE.

The following lines explain the resampling protocol that we have used. All the numbers used can easily be tuned in the scripts provided.

To test our models we split the data into two different datasets, one for training and one for testing. The testing dataset will contain $\frac{1}{3}$ of all the patients. As they are chosen randomly, it is expected that the proportion of patients in each of the classes is kept. We will only use the testing dataset to test our models.

In the training dataset we have 299 patients, and it is expected that about 254 of them will have the variable DIED = FALSE and 45 will have DIED = TRUE. As most of the models are sensitive to imbalanced datasets, we need to improve the balance of the training dataset. We use the *bagging* algorithm to generate 11 balanced datasets, each of them will be used to train a model. We chose the number of bags to be odd to avoid ties. We also tried different quantities of bags to the result that the F2 score does not improve any more with a number of bags larger than 11, so we determined this to be the quantity of bags to use. Then, to generate a prediction, each of the models will be used to make a *hard* vote, and the class predicted by the majority will be the result for an instance. This way we have constructed what we have called a *super-model* consisting of a list of 11 models that return a final vote for every instance of the testing dataset.

Each of the *bags* will contain the same amount of instances. As we want them to be balanced, all of them will contain every DIED = TRUE instance in the training dataset, and a random sample of the same size of the DIED = FALSE instances. Thus, each bag will have about $2 \times 45 = 90$ instances.

5 Metric used to evaluate the models

To evaluate our models we chose the F2 score, which is the harmonic mean of recall and precision with more emphasis on precision. We chose this score because the two result classes of our dataset are quite imbalanced. Because we are especially interested in a low number of false negatives, i.e. predictions that a patient will not die within a year when he actually does, we defined precision and recall as follows:

$$\begin{aligned}\text{precision} &= \frac{\text{True Negative}}{\text{True Negative} + \text{False Negative}} \\ \text{recall} &= \frac{\text{True Negative}}{\text{True Negative} + \text{False Positive}}\end{aligned}$$

The F2 score is then defined as:

$$F2 = (1 + 2^2) \cdot \frac{\text{precision} \cdot \text{recall}}{2^2 \cdot \text{precision} + \text{recall}}$$

One advantage of this score is that we can put more weight to the under-represented class. At the same time the F2 score has the disadvantage that it does not give a clear result such as for example the accuracy score because the F2 score consists of the combination of two different values leading to possibly similar F2 scores with different prediction results. Therefore we will not only trust blindly in the calculated F2 score but will consider the confusion matrices of each model as well to determine the best possible model.

6 Results obtained using linear/quadratic methods

6.1 Naive Bayes

The Naive Bayes Algorithm has the advantage that it does not distinguish between types of data and it does not perform any implicit transformation. The typical disadvantage of this method is that it assumes independence of the attributes of the data.

As we have mixed data, the advantage is very appropriate. Looking at our data, it seems like most of the attributes are independent, so we can assume the results will be good. In fact, as we are working with a reduced dataset (in which we have removed some of the attributes), we could expect that the results will be even better in that one.

As this is a quite simple method, it does not need parameters, so we do not have to choose any hyper-parameter.

The results obtained using this model are:

At the first sight, it stands out the fact that the results from the “original” dataset and the “removed” one are very different. While “original” only shows a F2 score about 0.3, the other one shows much better results with a F2 score of 0.94, as expected. This leads us to think that the attributes we removed weren’t indeed independent from the others, and keeping them out has helped the model.

Table 1: NB Original

Prediction	Reference	
	False	True
False	32	3
True	93	21

Table 2: NB Removed

Prediction	Reference	
	False	True
False	122	18
True	5	4

Table 3: F2 on NB

	Original	Removed
Naive Bayes	0.2990654	0.9413580

But also we can see that the model predicts the DIED = TRUE cases quite well in the "original" dataset, but predicts DIED = FALSE cases really badly. In the "removed" dataset it is vice versa. So for the Naive Bayes it seems not possible to get a good prediction rate for both classes.

6.2 K-nearest Neighbours

K-nearest Neighbours (KNN) will look for similar patients and will predict the the major class among them. This method is quite sensitive to very imbalanced data, since with many more instances in one class the probability of having these instances close to the test cases is very high. Hence, we hope that our resampling method will be suitable for this model, as it only trains it with a balanced dataset of patients.

KNN is also sensitive to non-normalized data. As stated earlier we did some normalization to our continuous variables, but also the used R method should apply further normalization if needed.

As most of our data is boolean, we have chosen to measure with the Manhattan distance instead of the Euclidean one. We expect results will be better with Manhattan distance since it reflects better the differences among patients.

Looking at the models created with crossvalidation we see that the average number of neighbours considered is 6.45 for "Original" and 5.91 for "Removed", which seems like normal values for k . The results obtained using this model are:

Table 4: KNN Original

Prediction	Reference	
	False	True
False	58	8
True	67	16

Table 5: KNN Removed

Prediction	Reference	
	False	True
False	73	8
True	54	14

Table 6: F2 on KNN

	Original	Removed
K-nearest neighbours	0.5123675	0.6196944

The $F2$ score is not very good for none of the datasets, but looking at the confusion matrices it seems that the model achieves very well the task of avoiding false negatives. This may happen because the $F2$ score is not the most suitable value for scoring the results, but it performs well for model parameters tuning.

6.3 General Linear Model

The main problem of using a General Linear Model (GLM) is that it doesn't distinguish between different data types. For this model all the values are real numbers. As most of our variables are categorical, we assume it will not reflect the data very well. In addition, the power of this method relies on the basis function that the user can define, according to the meaning he knows they may have. We have not defined any due to lack of medical knowledge, so this feature won't be used.

The results of running this model on the dataset are:

Table 7: GLM Original

Prediction	Reference	
	False	True
False	80	14
True	45	10

Table 8: GLM Removed

Prediction	Reference	
	False	True
False	86	13
True	41	9

Table 9: F2 on GLM

	Original	Removed
Linear Model	0.6734007	0.7084020

The $F2$ score is not as bad as expected. In fact, it performs better than KNN despite the number of false negatives being higher.

7 Results obtained using non-linear methods

7.1 Random Forest

Random forests build decision trees using bagging to get different samples and combinations of variables to train on, and classify given test data by majority vote of all decision trees. This method is said to be quite robust even to imbalanced data, but does not create bags with equal parts of both classes, so we chose to call the random forest method with our preprocessed data. As the already balanced train samples are then divided into bags the models should give a more or less good result in the testing phase. Although it is possible to define the depth of trees and number of variables to choose we did not set these parameters fixed, because after trying different combinations there was none to create results clearly superior to the others.

Table 10: RF Original

Prediction	Reference	
	False	True
False	76	11
True	49	13

Table 11: RF Removed

Prediction	Reference	
	False	True
False	78	12
True	49	10

Table 12: F2 on RF

	Original	Removed
Random Forest	0.6473595	0.6521739

The two confusion matrices show very similar results and thereby underline the characteristic of being robust, as our preprocessing does not significantly improve the results. This also shows in the $F2$ score of the two test samples which is 0.647 for the original dataset and with 0.652 only a little better for the processed dataset. For the parameters of the models the *caret* package has chosen only one tree and one variable to classify.

7.2 Neural Network

Neural Networks make use of the biological system of neurons nesting functions into each other to classify the given input data. In our case we used a single-hidden-layer neural network. We did not allow skipping because a test with

the skipping option set did not improve the results. We neither set a fix regularization parameter because this also let to poorer results when testing new data on the calculated models.

Table 13: NN Original

Prediction	Reference	
	False	True
False	56	10
True	69	14

Table 14: NN Removed

Prediction	Reference	
	False	True
False	60	6
True	67	16

Table 15: F2 on NN

	Original	Removed
Neural Net	0.4946996	0.5226481

The confusion matrices of the Neural Network also show an obvious similarity. It is remarkable that in both cases the neural network predicted about $\frac{3}{4}$ of the true instances correctly which is a good result. But at the same time the predictions of the false instances not even get 50% correctly. This leads to a lower F2 score which is 0.495 for the original and 0.523 for the processed dataset.

As parameters mostly a decay of $\lambda = 0$ and a single neuron are chosen, with few outliers of up to 5 neurons and a decay of $\lambda = 0.0001$.

8 Description and justification of the final model chosen and estimation of the generalization error

As we used the F2 score to evaluate the different models, it would be appropriate also to use it to choose the best model amongst them. The model with the highest F2 score is Naive Bayes, but only for one dataset while it at the same time has the lowest F2 score for the other dataset. The chosen model with the highest average F2 score ist the GLM which reaches a mean of 0.69.

The GLM also has an average prediction rate of 62%, calculated with the testing dataset, and therefore a generalization error of 38%. This is no really a desirable testing error rate, and even more so as the GLM predicts many false negatives which we tried to avoid. These results may be explained by different aspects which are explained in section 11.

9 Self-assessment of successes, failures and doubts

9.1 Successes

- Learn how to train and evaluate different models in a context
- Implementation of F2 score
- Improvement of R programming skills

9.2 Failures

- Found no model that predicts better than the "always false model" (evaluated with F2 score)
- Even though the scores of the models are not that bad, the confusions matrices show that in total many instances are not predicted correctly

9.3 Doubts

- It seems that for KNN the Euclidean metric works better than the Manhattan metric (evaluated with F2 score)
- It might be that there could be a metric more appropriate than the F2 score. The related paper states that the geometric mean was used in their work.

10 Scientific and personal conclusions

It seems appropriate that the results of the chosen model do not match our expectations of good results, because real problems can be quite complex and need more complex and advanced modelling to hopefully achieve desirable solutions. This work also shows the difficulty of solving medical questions, which in addition to being complex also cannot provide the desired size of datasets to improve the training of models.

As for our personal conclusion it can be said that this work gave us a more comprehensive view of Machine Learning problems and the algorithms trying to solve them. We got a small insight in what the work in this sector might be like and which problems and challenges there are to be faced. To conclude we can say that we learned a lot about the different models and algorithms and that this work strongly supported us in the development of a basic knowledge of Machine Learning.

11 Possible extensions and known limitations

Regarding the results of our chosen model we can see that even this model's rate of correct predictions is not very high. As we have only tried to predict with five different models, it is possible that a model we did not try to use would give better predictions for the given dataset. As the related paper suggests to use Support Vector Machines it might well be that such a model gives better predictions. Furthermore we have just begun learning how to use method of

Machine Learning and therefore have few experience, so we even might not have found the best parameters to use on the models we chose.

Also we do not have much medical knowledge of the data we have been working with, so we could neither put more emphasis on possibly more significant variables nor declare variables or values as not significant. The modifications we made to the dataset were only based on removing obviously underrepresented or unrealistic data. In cooperation with someone who has medical knowledge of the dataset's variables maybe the process of training could be improved.

As a last limitation the dataset only contains few instances to work with, which are even fewer after the elimination of some instances that we defined as outliers. As the dataset was published in 2013 by now there could be more data available to train on, which we assume would also lead to better results.

References

- [1] Nuvoair. *Do you know how to interpret the results of your spirometry test?* URL: <https://www.nuvoair.com/blog/do-you-know-how-to-interpret-the-results-of-your-spirometry-test.html>.
- [2] Maciej Zikeba et al. "Boosted SVM for extracting rules from imbalanced data in application to prediction of the post-operative life expectancy in the lung cancer patients". In: *Applied Soft Computing* (2013). DOI: [WebLink].